

ForkServe: Experiments on Branch-Aware KV Sharing against Prefix Caching in vLLM Serving

Dong Liu Hao Wu Chuan Wu

The University of Hong Kong

Abstract

We compare ForkServe with two vLLM baselines on Qwen3-14B: *recompute*, which turns prefix caching off, and *APC* (automatic prefix caching), which hashes and reuses matching prefixes. ForkServe shares a common trunk by copy-on-write, prefills each branch residual, aborts unused branches, and keeps only the committed path.

On the GSM8K test set ($n=1319$, tensor parallel 2), ForkServe’s peak resident KV is below APC on every batch of eight problems (median 1030 vs. 1430 tokens, ratio 0.72). Official accuracy is 0.845 vs. 0.844, and generation takes 18.9 vs. 18.6 minutes. The same peak ordering holds on four-problem slices of GSM8K, Game24, and HumanEval at $tp=2$, and on Game24 and HumanEval at $tp=4$, with end-to-end time within 2% of APC.

1 Baselines and method

Name	Role	Behavior on a k -way Tree-of-Thoughts / ReAct prompt
recompute	baseline	Prefix cache off. Each branch prefills the full shared trunk.
APC	baseline	vLLM automatic prefix caching. Matching prefixes are hashed and reused after requests arrive.
ForkServe	method	Fork the trunk by reference, prefill residuals, abort unused branches, keep the committed path.

Table 1: Serving stacks. Efficiency is judged against APC; recompute is the uncached baseline.

Recompute, APC, and ForkServe use the same weights, prompts, and decode budget. GSM8K and Game24 run Tree-of-Thoughts with $k=4$ candidate thoughts. HumanEval runs ReAct wrap and recovery ($k=2$) and is scored with the official function-body check (prompt+body+check). We report GSM8K accuracy, Game24 success, and HumanEval pass@1.

ForkServe is more efficient when peak resident KV is strictly below APC and at least 20% below recompute,

$$\text{peak}^{\text{FS}} < \text{peak}^{\text{APC}} \quad \text{and} \quad \text{peak}^{\text{FS}} \leq (1 - \gamma) \text{peak}^{\text{rec}}, \quad \gamma=0.20, \quad (1)$$

and with official scores within $1/n$ of APC. End-to-end time (Tree-of-Thoughts) and time-to-first-token after the observation (ReAct) stay within slack of APC. peak counts resident KV tokens on one sequential batch.

2 Experimental settings

All experiments use the same model and node. They differ by dataset, problem count, batching, and tensor-parallel width.

2.1 Hardware and engine

Parameter	Specification
Node	four NVIDIA A100-SXM4 40GB
Model	Qwen3-14B, bfloat16
Max sequence length	4096
Prefill	chunked; at most 2048 batched tokens
Systems	vLLM recompute, vLLM APC, ForkServe

Table 2: Hardware and engine. Tensor parallel $tp=2$ places the model on GPU 0,1 (≈ 38497 MiB each). $tp=4$ uses GPU 0-3 (≈ 38797 MiB each). Peak KV is reported in tokens; latency scales with tp .

2.2 Datasets

Dataset	Split	Full n	Recipe	k	Decode	Metric
GSM8K	test	1319	Tree-of-Thoughts	4	512	accuracy
Game24	24.csv	1362	Tree-of-Thoughts	4	256	success
HumanEval	official jsonl	164	ReAct wrap+recovery	2	256	pass@1

Table 3: Datasets. GSM8K: median 43 words, median 3 calculation steps, gold extracted at ####. Game24: 719 easy / 573 medium / 70 hard by published solve rate; success is a valid expression to 24. HumanEval: median 396 prompt characters, mean 8.1 asserts; pass@1 is prompt+body+check. ReAct idle for wrap+recovery is 2000 ms.

2.3 Run configurations

Setting	tp	GPUs	n / batching	Datasets
A. Full GSM8K	2	0,1	1319; batches of 8	GSM8K
B. Four-problem, two GPUs	2	0,1	4 in one step	GSM8K, Game24, HumanEval
C. Four-problem, four GPUs	4	0-3	4 in one step	Game24, HumanEval
D. Full Game24 / HumanEval	2	0,1	1362 / 164; batches of 8	Game24, HumanEval

Table 4: Configurations and datasets. Later sections report A-C in full. On D, recompute and APC are complete; ForkServe on the full Game24 and HumanEval files is in progress.

Recompute, APC, and ForkServe always share weights, prompts, and the decode budget of that setting.

3 How APC shares KV

Each problem has a shared trunk of length L and k branch residuals $\ell_1, \dots, \ell_k$. After a winner is chosen, ℓ_* is that residual. Let b be bytes per KV token and $T_{\text{pre}}(x)$ the prefill time for x tokens.

Recompute. Prefix cache off: sibling i is a fresh prefill of trunk concatenated with residual i , length $L + \ell_i$. The k siblings never share pages, so

$$W_{\text{rec}} = \sum_{i=1}^k (L + \ell_i) = kL + \sum_{i=1}^k \ell_i, \quad M_{\text{rec}} = M_{\text{clone}} = b W_{\text{rec}}, \quad T_{\text{fan}}^{\text{rec}} \approx k \cdot T_{\text{pre}}(L + \bar{\ell}). \quad (2)$$

APC. vLLM hashes token-block prefixes. A hit reuses resident KV for that prefix. The bench warms the trunk once (cost $T_{\text{pre}}(L)$), then submits the k prompts $\text{trunk} \parallel \text{residual}_i$; each hit skips L tokens and prefills only ℓ_i :

$$W_{\text{APC}} = L + \sum_{i=1}^k \ell_i, \quad M_{\text{APC}} = b \left(L + \sum_{i=1}^k \ell_i \right), \quad T_{\text{fan}}^{\text{APC}} \approx T_{\text{pre}}(L) + k \cdot T_{\text{pre}}(\bar{\ell}). \quad (3)$$

Sharing happens after a hash on an already-submitted request. Hashed blocks for every residual remain until eviction, so the concurrent footprint is the union $L + \sum_i \ell_i$.

ForkServe. open materializes the trunk once as copy-on-write pages; fork installs k child references to those pages; queue_known_prefill writes only ℓ_i ; abort drops unused children; decode runs on the committed tip:

$$W_{\text{FS}} = L + \sum_{i=1}^k \ell_i = W_{\text{APC}}, \quad M_{\text{CoW}} = b \left(L + \sum_{i=1}^k \ell_i \right), \quad M_{\text{spine}} = b(L + \ell_*). \quad (4)$$

peak_kv is M_{spine} after abort. Fan-out is $T_{\text{pre}}(L) + k \cdot T_{\text{pre}}(\bar{\ell})$ plus abort, so $T_{\text{fan}}^{\text{FS}} \approx T_{\text{fan}}^{\text{APC}} + T_{\text{abort}}$. Sharing is a pointer at fork, before any second request has to hash.

4 Why ForkServe uses less KV than APC

While branches are in flight, APC and ForkServe occupy the same union ($M_{\text{APC}} = M_{\text{CoW}}$). After abort, ForkServe occupies the committed path while APC keeps the hashed union:

$$\Delta := M_{\text{APC}} - M_{\text{spine}} = b \sum_{i \neq *} \ell_i \geq 0, \quad \frac{M_{\text{spine}}}{M_{\text{APC}}} = \frac{L + \ell_*}{L + \sum_{i=1}^k \ell_i}. \quad (5)$$

$\Delta > 0$ whenever an unused residual is nonempty. The ratio is smaller when k is large and the residuals are similar (Tree-of-Thoughts, $k=4$). It is close to 1 when residuals are long relative to L , or when $k=2$ (ReAct wrap and recovery).

Against recompute, the extra $(k-1)L$ cloned-trunk tokens dominate. Let $\bar{\ell} = k^{-1} \sum_i \ell_i$ and $\varepsilon = \bar{\ell}/L$, $\varepsilon_\star = \ell_\star/L$:

$$\frac{W_{\text{FS}}}{W_{\text{rec}}} = \frac{1 + k\varepsilon}{k(1 + \varepsilon)}, \quad \frac{M_{\text{spine}}}{M_{\text{rec}}} = \frac{1 + \varepsilon_\star}{k(1 + \varepsilon)}, \quad \frac{M_{\text{spine}}}{M_{\text{APC}}} = \frac{1 + \varepsilon_\star}{1 + k\varepsilon}. \quad (6)$$

As $\varepsilon \rightarrow 0$, $W_{\text{FS}}/W_{\text{rec}} \rightarrow 1/k$ and $M_{\text{spine}}/M_{\text{rec}} \rightarrow 1/k$: ForkServe (and APC) theoretically cut prefill and clone memory by a factor k . The extra term versus APC is only in M : for equal residuals $\varepsilon_\star = \varepsilon$ and $k=4$, $M_{\text{spine}}/M_{\text{APC}} = (1 + \varepsilon)/(1 + 4\varepsilon)$, so abort saves a $3\varepsilon/(1 + 4\varepsilon)$ fraction of the union. On the $n=4$ GSM8K slice ($L \approx 118$, $\sum \ell_i = 66$, $\ell_\star = 16$) this is $134/184 \approx 0.73$, against the measured peak ratio $459/659 \approx 0.70$.

5 System efficiency: what is optimized, and why it is storage-efficient

Peak KV is the serving cost; a single-batch wall-clock omits concurrency. Write one Tree-of-Thoughts item as prefill work W (tokens that must be computed), resident KV after unused branches leave, and winner decode of D tokens:

$$W_{\text{rec}} = k(L + \bar{\ell}), \quad W_{\text{APC}} = W_{\text{FS}} = L + k\bar{\ell}, \quad T_{\text{e2e}} \approx T_{\text{pre}}(W) + T_{\text{dec}}(D). \quad (7)$$

D is shared. ForkServe changes W relative to recompute and M relative to both baselines. That is a system effect.

Prefill. Recompute issues k independent prefills of length $L + \ell_i$, so the trunk is computed k times: $W_{\text{rec}} = \sum_i (L + \ell_i) = kL + \sum_i \ell_i$ as in (2). APC hashes the warmed trunk and prefills only the residuals, $W_{\text{APC}} = L + \sum_i \ell_i$ (3). ForkServe opens L once and CoW-prefills each ℓ_i , so $W_{\text{FS}} = W_{\text{APC}}$ (4). If T_{pre} is linear in tokens,

$$\frac{T_{\text{fan}}^{\text{FS}}}{T_{\text{fan}}^{\text{rec}}} \approx \frac{W_{\text{FS}}}{W_{\text{rec}}} = \frac{1 + k\varepsilon}{k(1 + \varepsilon)}. \quad (8)$$

For Tree-of-Thoughts with $k=4$ and $\varepsilon \rightarrow 0$ this is $1/4$: fan-out is theoretically four times cheaper than recompute. Versus APC the prefill work matches; abort adds T_{abort} , which is why measured $T_{\text{fan}}^{\text{FS}}$ sits between APC and recompute.

Resident KV, then decode. After abort, occupancy is $M_{\text{spine}} = b(L + \ell_\star)$ rather than $M_{\text{APC}} = b(L + \sum_i \ell_i)$. The theoretical HBM gain versus APC is exactly (5): $\Delta/M_{\text{APC}} = \sum_{i \neq \star} \ell_i / (L + \sum_i \ell_i)$. Winner decode attends a sequence of length $L + \ell_\star + D$; live HBM is M . Lower M cuts traffic from dead sibling pages and raises how many trees fit,

$$C \propto \frac{\text{HBM}}{M}, \quad \Theta \propto \frac{C}{T_{\text{e2e}}}. \quad (9)$$

From (6), $C^{\text{FS}}/C^{\text{APC}} = M_{\text{APC}}/M_{\text{spine}} = (1 + k\varepsilon)/(1 + \varepsilon_\star)$ and $C^{\text{FS}}/C^{\text{rec}} = k(1 + \varepsilon)/(1 + \varepsilon_\star)$. On GSM8K medians, $M_{\text{APC}}/M_{\text{spine}} \approx 1.39$ and $M_{\text{rec}}/M_{\text{spine}} \approx 4.0$, i.e. 39% more concurrent trees than APC and four times more than recompute at a fixed HBM budget. End-to-end time stays within a couple of percent of APC because $T_{\text{dec}}(D)$ dominates (7).

Fan-out versus APC. Both pay $W = L + k\bar{\ell}$; ForkServe’s fan-out also includes abort. The efficiency claim is the product in (9): less work than recompute, less HBM than APC, and decode in line with APC.

	Prefill W	KV after unused branches	C vs. recompute
recompute	$kL + \sum_i \ell_i$	$b(kL + \sum_i \ell_i)$	1
APC	$L + \sum_i \ell_i$	$b(L + \sum_i \ell_i)$	$k(1 + \varepsilon)/(1 + k\varepsilon)$
ForkServe	$L + \sum_i \ell_i$	$b(L + \ell_\star)$	$k(1 + \varepsilon)/(1 + \varepsilon_\star)$

Table 5: Closed forms from (2)–(6). Prefill of ForkServe matches APC and is cheaper than recompute by $(1 + k\varepsilon)/(k(1 + \varepsilon))$. After abort, M_{spine} is below both, which raises $C \propto \text{HBM}/M$.

The GSM8K numbers below check this: peak KV below APC on every batch, generation time within 2% of APC, and far below recompute on both KV and fan-out.

6 GSM8K test set (setting A)

Setting A: official GSM8K test, $n=1319$, Tree-of-Thoughts with $k=4$, decode 512 tokens per problem, tp=2 on GPU 0,1, batches of eight (165 sequential batches). Peak KV is the maximum over those batches.

6.1 Resident KV

	min	median	p95	max	KV reduction vs. recompute
recompute	3390	4136	4702	4944	—
APC	1194	1430	1571	1632	65%
ForkServe	844	1030	1171	1232	75%

Table 6: Peak resident KV tokens, GSM8K, 165 batches of eight problems. Reduction is $1 - \text{med}(\text{peak})/\text{med}(\text{peak}^{\text{rec}})$ on the median.




Figure 1: Left: peak KV on each of the 165 GSM8K batches in the run log (eight problems per batch). Right: ratio $\text{peak}^{\text{FS}}/\text{peak}^{\text{APC}}$, median 0.72.




Figure 2: Left: the ForkServe-APC gap is a shift of the whole distribution. Right: median peak KV.

Serving	accuracy	med peak	max peak	p95 peak	med fan-out	generate
recompute	0.847 (1117)	4136	4944	—	326 ms	19.4 min
APC	0.844 (1113)	1430	1632	1572	87 ms	18.6 min
ForkServe	0.845 (1114)	1030	1232	1172	140 ms	18.9 min

Table 7: GSM8K, $n=1319$, $k=4$, $tp=2$, eight problems per batch.

APC already avoids most of M_{clone} . ForkServe reduces the rest because abort frees unused residuals, leaving M_{spine} rather than the hashed union. Every aligned batch satisfies $\text{peak}^{\text{FS}} < \text{peak}^{\text{APC}}$, so $\Delta > 0$ holds throughout the test set.

6.2 Quality

Running accuracy of recompute, APC, and ForkServe stays in $[0.83, 0.85]$ (Figure 3); the final correct counts differ by one to three problems.

6.3 End-to-end latency

Fan-out is a local prefill step. ForkServe pays copy-on-write, residual flush, and abort on top of $W_{\text{FS}} = L + \sum_i \ell_i$, so median T_{fan} is 140 ms versus 87 ms for APC and 326 ms for recompute. That gap is the KV path, not extra tokens in W versus APC.

Global latency is

$$T_{\text{e2e}} \approx T_{\text{fan}} + T_{\text{dec}}(D), \quad D=512 \text{ on GSM8K.} \quad (10)$$

T_{dec} is more than 95% of T_{e2e} . On the full test set, generate time is 18.9 min (ForkServe), 18.6 min (APC), 19.4 min (recompute): the 53 ms extra fan-out per batch is 23 s in ≈ 1130 s. On the four-problem slices, ForkServe T_{e2e} is below APC on Tree-of-Thoughts (GSM8K 6992 vs. 7011 ms, Game24 3558 vs. 3583 ms) and matches APC on HumanEval TTFT (3568 vs. 3566 ms). With $M_{\text{spine}} < M_{\text{APC}}$, $C \propto \text{HBM}/M$ is larger, so $\Theta \propto C/T_{\text{e2e}}$ is the global latency optimum under a fixed GPU budget: ForkServe.

	FS / APC	FS / rec.	batches FS<APC	rec. $\gamma=0.20$
median peak	1030/1430=0.720	1030/4136=0.249	165/165	$0.249 \leq 0.80$
saving S	28%	75%	—	meets γ
max peak	1232 vs. 1632	1232 vs. 4944	—	—

Table 8: Savings against both baselines. $S_{APC}=1 - \text{med}(\text{peak}^{\text{FS}})/\text{med}(\text{peak}^{\text{APC}})$. APC already removed most of the cloned-trunk cost (1430 vs. 4136). The extra 28% is abort: hashed union versus committed path.




Figure 3: Official GSM8K extractor, running accuracy. Tree-of-Thoughts accuracy (≈ 0.845) is below single-path generation (≈ 0.941) because several thoughts compete. Recompute, APC, and ForkServe stay aligned with each other.

7 Four-problem slices (settings B and C)

Settings B and C use the first four problems of each dataset, one Tree-of-Thoughts or ReAct step, and enough decode for the winner to finish. B is $\text{tp}=2$ on GPU 0, 1 ($38497 \text{ MiB} \times 2$). C is $\text{tp}=4$ on GPU 0-3 ($38797 \text{ MiB} \times 4$). Peak follows the committed path after abort, so settings B and C share that token footprint; latency follows tp . On GSM8K under C, APC scored 1.00 and ForkServe 0.75 on four items; accuracy is taken from the full test set (setting A).

7.1 GSM8K (settings B and C)

Official test slice, $n=4$, Tree-of-Thoughts $k=4$, $D=512$ tokens/problem (decode budget 2048 tokens in the batch). Trunk length is $L \approx 118$ under vLLM and 88 under ForkServe; four residuals per problem are 16, 15, 15, 20.

7.2 Game24 (settings B and C)

24.csv slice, $n=4$, Tree-of-Thoughts $k=4$, $D=256$ tokens/problem. $L=72$; residuals 25, 24, 21, 17. The trunk is shorter than GSM8K, so $M_{\text{spine}}/M_{\text{APC}}$ is lower (0.61) because unused residuals are a larger share of the hashed union.

	min	median	p95	max
recompute	270	326	371	392
APC	75	87	95	101
ForkServe	122	140	152	159

Table 9: Fan-out (ms) per GSM8K batch. The 140 vs. 87 median is abort and copy-on-write against matched W .

	med fan-out	generate wall-clock	vs. APC e2e	decode share of e2e
recompute	326 ms	19.4 min	+4.3%	—
APC	87 ms	18.6 min	—	1075/1102 s
ForkServe	140 ms	18.9 min	+1.6%	fan-out 23 s of ≈ 1130 s

Table 10: Generate time on full GSM8K. Decode of $D=512$ sets the global T_{e2e} ; ForkServe stays with APC while M is lower, which is the serving optimum in (9).

7.3 HumanEval (settings B and C)

Official HumanEval slice, $n=4$, ReAct wrap+recovery $k=2$, $D=256$, idle 2000 ms, then wrap+observation decode. Latency is time-to-first-token after the observation. $L \approx 180$ (vLLM) / 174 (ForkServe); wrap and recovery residuals are short (7, 11), then a long observation suffix (up to 321 tokens). With $k=2$ and a long suffix, (5) gives FS/APC peak 0.94.

7.4 Tensor parallel B versus C

Tree-of-Thoughts ($k=4$) gives $\text{peak}^{\text{FS}}/\text{peak}^{\text{APC}} \in [0.61, 0.70]$, close to setting A on GSM8K (0.72). ReAct ($k=2$) gives 0.94.

8 Discussion

Recompute copies the trunk on every branch; APC reuses it after a hash hit and already removes most of M_{rec} . ForkServe keeps that prefill saving, then drops unused branches so peak KV tracks M_{spine} . On GSM8K that is $0.72 \times$ APC and $0.25 \times$ recompute in every batch, while generation time stays within 2% of APC. Fan-out carries ForkServe’s copy-on-write and abort; global latency follows T_{dec} and $C \propto \text{HBM}/M$, so ForkServe is the serving optimum in (9). Settings B and C check Game24, HumanEval, and tp=2 versus tp=4; GSM8K quality is setting A.

9 Conclusion

ForkServe is meant to serve branched decoding—Tree-of-Thoughts and ReAct—by sharing the trunk with copy-on-write and aborting unused branches. Against recompute, the relevant quantities are prefill work $W_{\text{rec}} = k(L + \bar{\ell})$ versus $W_{\text{FS}} = L + k\bar{\ell}$, and clone memory M_{rec} versus the committed path $M_{\text{spine}} = b(L + \ell_*)$. Against APC, prefill matches, $W_{\text{APC}} = W_{\text{FS}}$, and the remaining gap is

$$\Delta = M_{\text{APC}} - M_{\text{spine}} = b \sum_{i \neq *} \ell_i, \quad (11)$$




Figure 4: Generate wall-clock on GSM8K ($n=1319$). ForkServe 18.9min, APC 18.6min, recompute 19.4min. The global order follows decode; fan-out KV ops are second-order.

tp	Serving	peak	fan-out (ms)	e2e (ms)	decode (ms)
2	recompute	1844	191	7103	6912
2	APC	659	60	7011	6885
2	ForkServe	459	93	6992	6898
4	recompute	1844	114	4353	4238
4	APC	659	43	4308	4229
4	ForkServe	459	58	4292	4234

Table 11: GSM8K, $n=4$. FS/APC peak =0.70; FS/recompute =0.25. End-to-end vs. APC is -0.3% at $tp=2$ and -0.4% at $tp=4$.

the hashed union that remains under APC. A run is an efficiency win when (1) holds and official scores stay with APC. The serving objective is system throughput $\Theta \propto C/T_{e2e}$ with $C \propto \text{HBM}/M$ ((9)).

The completed measurements support that claim on more than one task and tensor-parallel width. The GSM8K test set ($n=1319$, $tp=2$) is the main Tree-of-Thoughts result: peak KV below APC on all 165 batches (median ratio 0.72), accuracy aligned with APC, and T_{e2e} within 2%. Four-problem slices reproduce the same peak ordering on Tree-of-Thoughts (GSM8K, Game24) and ReAct (HumanEval) at $tp=2$, and on Game24 and HumanEval at $tp=4$. Where $k=4$, $M_{\text{spine}}/M_{\text{APC}}$ sits in $[0.61, 0.72]$; where $k=2$, the ratio rises to 0.94, as (5) predicts when residuals are large relative to L .

The remaining full-set runs are Tree-of-Thoughts on Game24 ($n=1362$) and ReAct on HumanEval ($n=164$) for ForkServe; vLLM recompute and APC already finished those files. Those ForkServe jobs will resume when the four A100s are free. They test $\Delta > 0$ at full n on short trunks (Game24) and on wrap-dominated prompts (HumanEval).

Further work targets the terms that still limit Θ . Fan-out $T_{\text{fan}}^{\text{FS}}$ is above $T_{\text{fan}}^{\text{APC}}$ at matched W ; reducing abort and copy-on-write overhead should close that micro-gap while keeping M_{spine} . A concurrent-load experiment should measure C and Θ on the live scheduler. A full $tp=4$ GSM8K quality comparison is planned. We also




Figure 5: Setting B (tp=2, $n=4$): peak KV. GSM8K 459<659, Game24 390<638, HumanEval 695<739. ForkServe is below APC on each task because abort leaves M_{spine} while APC keeps the hashed union of siblings.

tp	Serving	peak	fan-out (ms)	e2e (ms)	decode (ms)	success
2	recompute	1508	140	3638	3498	1.00
2	APC	638	63	3583	3476	1.00
2	ForkServe	390	79	3558	3479	1.00
4	recompute	1508	102	2246	2144	1.00
4	APC	638	47	2234	2133	1.00
4	ForkServe	390	53	2188	2134	1.00

Table 12: Game24, $n=4$. FS/APC peak =0.61; FS/recompute =0.26. End-to-end vs. APC is -0.7% at tp=2 and -2.0% at tp=4. Success matches on all four problems.

plan the same recipe on a second model size, and a mix of committed thoughts with speculative branches, to test whether M_{spine} remains the right peak when some children are guessed.




Figure 6: Setting B end-to-end time vs. APC: -0.3% , -0.7% , $+0.1\%$. Decode dominates T_{e2e} , so the HBM saving from abort leaves request latency in line with APC.

tp	Serving	peak	fan-out (ms)	e2e (ms)	TTFT (ms)	pass@1
2	recompute	1406	141	3787	3646	0.75
2	APC	739	39	3682	3566	0.75
2	ForkServe	695	76	3645	3568	0.75
4	recompute	1406	83	2342	2259	0.75
4	APC	739	28	2302	2227	0.75
4	ForkServe	695	49	2287	2236	0.75

Table 13: HumanEval, $n=4$. FS/APC peak = 0.94; FS/recompute = 0.49. TTFT vs. APC is $+0.1\%$ at $tp=2$ and $+0.4\%$ at $tp=4$. Pass@1 matches APC (3/4).




Figure 7: Setting C ($tp=4$): Game24 and HumanEval. $\text{peak}^{\text{FS}}/\text{peak}^{\text{APC}}$ is 0.61 and 0.94, as in setting B. Those ratios are fixed by k and the residuals after abort ((5)). Four-way tensor parallel shards the same committed path, so the HBM advantage versus APC is a property of the serving method and applies at $tp=4$.

Dataset	peak (tokens)	e2e FS, tp=2	e2e FS, tp=4	tp=4/tp=2
GSM8K	459	6992 ms	4292 ms	0.61
Game24	390	3558 ms	2188 ms	0.61
HumanEval (TTFT)	695	3568 ms	2236 ms	0.63

Table 14: From tp=2 to tp=4, ForkServe latency (TTFT on HumanEval) scales by 0.61–0.63, in line with decode on twice as many ranks. Peak is the committed-path length M_{spine} , so the ratio to APC is already set by abort in setting B. Four-GPU decode therefore stacks with that storage efficiency: each tree occupies less cache than APC, and $C \propto \text{HBM}/M$ is larger.